

A Extended Supplementary Material

This appendix provides complete benchmark coverage, annotation details, prompt templates, and extended quantitative results. The main text reports compact results for readability, while this appendix preserves full technical details for reproducibility.

B Overview

In this supplementary material, we provide additional details on data construction, dataset statistics, implementation prompts, qualitative analysis, and benchmark examples.

C Data Construction and Annotation Details

C.1 Annotation Platform

To ensure annotation quality and professional consistency, we developed a custom annotation platform for fine-grained cinematic analysis. The tool integrates frame-level navigation, structured question answering, and reviewer correction.

- **Frame-level control.** Annotators can inspect clips frame-by-frame to identify transient cinematic patterns.
- **Taxonomy alignment.** The annotation panel follows the same fine-grained taxonomy used in the main paper.
- **Review mechanism.** Senior reviewers can audit and correct annotations in a dedicated pass.

D Additional Dataset Statistics

D.1 Duration Distribution Analysis

We analyze shot-duration statistics to characterize temporal complexity. The movie subset shows a long-tail distribution centered around short-to-medium durations, while different source films present distinct pacing styles. AI-generated clips are not included in this plot due to near-fixed generation duration in the tested setup.

E Evaluation Prompts and Settings

E.1 Prompt Template

We use a unified prompt format for all MLLMs and keep answer extraction deterministic.

E.2 Shot-boundary Preparation

Shot-boundary detection is used to generate candidate segments before expert review. Frames are resized and converted to grayscale, then compared by mean absolute frame difference. Candidate cuts are accepted only when a minimum temporal gap is satisfied and local difference exceeds an adaptive threshold. Nearby cuts are merged, and extreme spikes trigger secondary checks to avoid artifacts from flashes or post-production effects. Final boundaries are always confirmed by human reviewers.

E.3 Taxonomy and Prompt Structure

The CineBench taxonomy has 16 labels in four dimensions: Cinematography, Lighting, Color, and Emotional Cue. The comprehension prompt uses a fixed format across models, with model-specific wrappers only for input transport. The core prompt template is listed below.

System prompt. You are a professional cinematographer and film critic. Analyze the given video clip using cinematic-language terminology.

User prompt.

<Video frames>

Question: Which of the following shot scales appears in the video?

Options: (A) extreme long shot (B) medium shot (C) long shot (D) extreme close-up (E) distant shot

Answer with the option letter only.

E.4 Qualitative Example

F Benchmark Data Examples

F.1 Ethics, Access, and Governance

CineBench is intended for academic research on cinematic-language modeling. The release package includes metadata, labels, prompt templates, and evaluation code. For content with copyright restrictions, we provide controlled access instructions and index-level assets instead of unrestricted redistribution. The project documentation also reports known risks, including stylistic bias and possible misuse in synthetic media production.

F.2 Full Comprehension Results (16 and 64 Frames)

F.3 Full Generation Metrics

For reproducibility and page-limit control, we provide all seven prompt-group generation case studies, per-model outputs, and manual audit notes in the project repository linked in the abstract. The repository package follows the same evaluation protocol as the paper and includes failure tags for camera geometry, lighting direction, color tone, and emotion consistency.




Figure 5: Overview of the CineBench annotation platform.

Table 6: Extended sample entries from CineBench. Q-ID indicates question identifier.

Video ID	Q-ID	Category	Question	Options	Ans.
Source: AI Generated (Wan2.2)					
AI_1	AI_1_1	Cine.	Which of the following shot types appears in the video?	(A) single shot (B) fast-motion shot (C) semi-subjective shot (D) two shot (E) group shot	A
AI_1	AI_1_3	Cine.	Which of the following shot scales appears in the video?	(A) extreme long shot (B) medium shot (C) long shot (D) extreme close-up (E) distant shot	B
AI_1	AI_1_5	Light.	Which of the following best describes the nature of the primary light source?	(A) natural (direct light) (B) artificial light (C) natural (diffused light) (D) practical/ambient scene light	A
Source: Movie - Joker (JK)					
JK_1	JK_1_2	Cine.	Which of the following shot types appears in the video?	(A) establishing shot (B) single shot (C) objective shot (D) fast-motion shot (E) three shot	B
JK_1	JK_1_6	Cine.	Which of the following composition styles appears in the video?	(A) balanced composition (B) central composition (C) short-side composition (D) framed composition (E) negative space composition	A
Source: Movie - The Godfather (GF)					
GF_1	GF_1_1	Light.	Which of the following best describes the level of contrast (light vs. dark)?	(A) high contrast (B) low contrast	B
GF_1	GF_1_3	Emotion	Which of the following overall atmospheres or stylistic moods does the video convey?	(A) bizarre (B) everyday (C) tranquil (D) mystery (E) solemn	E




Figure 6: Temporal statistics of CineBench. Top: aggregate duration distribution. Bottom: per-film duration spread comparison.

Table 7: Evaluation results on CineBench. Dimension abbreviations: Cine (Cinematography), Light (Lighting), Color (Color), Emotion (Emotional Cue). Updated with latest models as of late 2025.

Model	frames	Avg	Cine	Light	Color	Emotion
<i>Closed-Sourced Models</i>						
GPT-5.2	16	65.1	66.0	66.3	60.0	62.0
	64	61.7	68.4	55.3	47.8	61.1
Gemini 3.0 Pro Preview	16	57.8	53.0	66.3	52.0	62.0
	64	56.8	61.6	55.3	42.2	52.2
Gemini 3.0 Flash Preview	16	58.9	60.0	63.8	46.0	52.0
	64	57.9	62.1	56.0	41.1	58.9
Gemini 2.5 Pro (think)[9]	16	53.5	55.0	60.0	37.5	42.5
	64	54.5	53.0	64.0	39.0	49.0
Gemini 2.5 Flash (think)[9]	16	47.2	45.0	58.0	20.4	53.0
	64	50.2	49.0	57.0	28.0	58.0
Doubao-Seed-1.8-vision-251228[7]	16	60.0	63.2	58.7	43.3	64.4
	64	60.6	64.7	60.0	40.0	62.2
Doubao-Seed-1.6-vision-250815 (think)[7]	16	60.1	62.1	66.0	37.5	55.0
	64	62.2	65.7	67.0	42.5	50.0
Qwen3-VL-Plus-2025-12-19[29]	16	59.9	60.5	62.0	47.8	62.2
	64	59.2	62.1	58.0	43.3	64.4
Qwen3-VL-Plus-2025-09-23 (think)[29]	16	50.6	50.0	54.0	39.0	55.0
	64	59.1	64.3	58.0	33.8	61.5
Qwen3-VL-Flash-2025-10-15 (think)[29]	16	51.4	52.0	57.0	32.3	51.0
	64	51.8	53.0	52.0	45.0	52.0
<i>Open-Sourced Models</i>						
Qwen3-VL-235B-A22B-Thinking[29]	16	56.2	59.5	56.0	35.6	61.1
	64	60.1	65.3	57.3	40.0	62.2
InternVL3-78B[41]	16	26.8	19.0	43.0	21.0	23.0
	64	28.8	21.0	48.0	25.0	14.0
Qwen3-VL-8B-Instruct[29]	16	49.3	46.0	56.0	31.1	64.0
	64	49.8	50.0	54.0	31.5	54.0
Qwen3-VL-8B-Thinking[29]	16	48.3	44.0	57.0	30.0	62.0
	64	51.8	53.0	56.0	32.3	53.0
GLM-4.5V-106B (think)[33]	16	54.5	57.6	51.0	50.0	54.1
<i>Human Baseline</i>						
Human	full	74.79	78.41	71.86	61.27	79.24